

## Solutions to exercises: week 2

(a) To prove that  $e^{-x} \geq (1-x)$  it is good to first make a sketch. You will see that only for  $x = 0$  we have equality. So first show that only for  $x = 0$  we have equality. Now you can derive that the derivative of  $e^{-x}$  is always larger (i.e. less negative) than the derivative of  $1-x$  for  $x > 0$ , which means that for  $x > 0$  will always hold  $e^{-x} \geq (1-x)$ . For  $x < 0$  the opposite holds (the derivative of  $e^{-x}$  is more negative). Now because if only for  $x = 0$  they are equal, then for  $x < 0$  we need to have that  $e^{-x} \geq (1-x)$ .

### Exercise 2.2

(a) To show the equivalence I will express the variables of the paper in terms of the variables of the slides. When the same variable is used (like  $y_i$ ) I will use a tilde to indicate variables from the paper.

1. Trivially, the indices are different:  $k, K \leftarrow t, T$
2. The labels are  $\{0, 1\}$  instead of  $\pm 1$ :  $\tilde{y}_i = \frac{1}{2}(y_i + 1)$ . This should also hold for the weak learners:  $h_t(\mathbf{x}) = \frac{1}{2}(f_k(\mathbf{x}) + 1)$ .
3. The error:  $\tilde{\epsilon}_t = \sum_i ((\tilde{w}_i^t / \sum_j \tilde{w}_j^t) |\frac{1}{2}f_k(\mathbf{x}_i) - \frac{1}{2}y_i|) = \sum_i ((c \cdot w_i^k / \sum_j c \cdot w_j^k) \mathcal{I}(f_k(\mathbf{x}_i) \neq y_i)) = \frac{\epsilon_k}{\sum_j w_j^k}$  where I defined  $\tilde{w}^t = c^k \cdot w^k$ .
4. Now it becomes more interesting. To rewrite  $\beta$ , it is actually easier to take the logarithm:  $\log \beta_t = \log(\frac{\epsilon_k / \sum_j w_j^k}{1 - \epsilon_k / \sum_j w_j^k}) = \log(\frac{\epsilon_k}{\sum_j w_j^k - \epsilon_k}) = -\log(\frac{\sum_j w_j^k - \epsilon_k}{\epsilon_k}) = -2\alpha_k$ , so therefore  $\beta_t = \exp(-2\alpha_k)$ .
5. To see that the weight update in the article (step 5) is equivalent as in the slides (step 4), note that  $1 - |h_t(\mathbf{x}_i) - y_i| = -y_i \alpha_k f_k(\mathbf{x}_i) - \alpha_k$ . Note that the article added weak learner  $t+1$ , while in the slides weak learner  $K$  is added. Then use that  $\beta^c = (\exp(-2\alpha))^c = \exp(-2c\alpha)$ , and that  $\exp(\sum_i c_i) = \prod_i \exp(c_i)$ , and see that these update rules are equivalent upto a factor  $\exp(\alpha_k)$ :  $\tilde{w}_i^{t+1} = \tilde{w}_i^t \exp(-y_i \alpha_K f_K(\mathbf{x}_i) - \alpha_K) = \prod_k \exp(-y_i \alpha_k f_k(\mathbf{x}_i)) \exp(-\alpha_k)$ . It seems that the constant  $c^k$  in step 3. was actually  $c^k = \exp(-\alpha_k)$ .

### Exercise 2.3

(a) My implementation is quite brute-force. The decision stump looks like:

```
function pred = decstump(x,bestf,bestt,bests)
```

```
if (bests>0)
    pred = (x(:,bestf)>=bestt);
else
    pred = (x(:,bestf)<bestt);
end
```

and the exhaustive search looks like:

```
function [bestfeat,besttheta,bestsign] = learndecstump(X,y,w)
p = w/sum(w);

besterr = inf;
bestfeat = 0;
besttheta = 0;
bestsign = +1;

[N,dim] = size(X);
% labels +-1
y = 2*y-1;
Nplus = sum(p(y>0));
```

```

Nmin = sum(p(y<0));

for i=1:dim
    [sx,I] = sort(X(:,i));
    sy = y(I); % sorted labels
    wsy = p(I).*sy; % weighted
    cy = cumsum([Nmin; wsy]);
    [mine,I] = min(cy);
    if (mine<besterr)
        besterr = mine;
        bestfeat = i;
        if I==1
            besttheta = sx(1)-10*eps;
        elseif (I>N)
            besttheta = sx(end)+10*eps;
        else
            besttheta = mean(sx((I-1):I));
        end
        bestsign = +1;
    end
    cy = cumsum([Nplus; -wsy]);
    [mine,I] = min(cy);
    if (mine<besterr)
        besterr = mine;
        bestfeat = i;
        if I==1
            besttheta = sx(1)-10*eps;
        elseif (I>N)
            besttheta = sx(end)+10*eps;
        else
            besttheta = mean(sx((I-1):I));
        end
        bestsign = -1;
    end
end
end

```

#### Exercise 2.4

(a) You should find that the best feature is 1, with a threshold around 1. Rescaling feature 2 should not have an influence.

#### Exercise 2.5

(a) The decision stump gave me an apparent error of around 0.18, but a test error of 0.51.

#### Exercise 2.6

Something like: `function [bestf,bestt,bests,beta] = learnadaboost(x,y,T)`

```

N = size(x,1);
w = ones(N,1)/N;
bestf = zeros(T,1);
bestt = zeros(T,1);
bests = zeros(T,1);
err = zeros(T,1);

for t=1:T
    p = w/sum(w);

```

```

[bestf(t),bestt(t),bests(t)] = learndecstump(x,y,p);
pred = decstump(x,bestf(t),bestt(t),bests(t));
df = abs(pred-y);
err(t) = df'*p;
if (err(t)==0)
    bestf = bestf(t,:);
    bestt = bestt(t,:);
    bests = bests(t,:);
    beta = 1/exp(1);
    break
end
beta(t) = err(t)/(1-err(t));
w = w.*(beta(t).^(1-df));
end

```

### Exercise 2.7


(a) When you test your implementation on a simple dataset, you should see that objects that are in a region of large class-overlap will obtain a high weight.

### Exercise 2.8

(a) The classification error on the Fashion dataset is around 0.29, way better than a decision stump.

### Exercise 2.9


(a) It is not so clear how the error depends on the number of iterations  $T$ . I got this graph:



It seems that after already 300 decision stumps, not so much improves.

### Exercise 2.10

(a) If you plot the learning curve for an Adaboost with  $T_{\max}=100$ :



you see that it converges for around  $N = 100$ . You also see that training errors are very small (basically 0 for **fashion57**!).