

# Computational Learning Theory

D.M.J. Tax




# Contents

- PAC learning, the definitions
  - Example: Rectangle Learning
  - Discrete hypothesis space and Consistent learners
  - Continuous hypothesis space: VC-dimension
  - 'No Free Lunch' theorem
- 
- Weak/strong learning
  - Boosting
  - AdaBoost

# PAC learning

- Probably Approximately Correct: PAC
- Here: restricted to boolean valued concepts from **noise-free** training data (often discrete features, although it can be extended...)
- Goal: learn a concept  $c$  from instances randomly drawn from prob.distribution  $D$  using learner  $L$ .




# PAC learning

- $X$ : instance space (all possible instances)
- $C$ : set of target concepts that may have to be learned
- $c$ : a concept, a subset of  $X$      $c : X \rightarrow \{0, 1\}$
- $D$ : probability distribution over instances  $x$ .
- $H$ : possible hypotheses used for approximating the concept  $c$     ( $H$  should include  $C$ )
- $L$ : learner that selects a hypothesis  $h$  given a random sample of instances drawn according to  $D$
- error:     $\text{error}_D(h) = \Pr_{x \in D} [c(x) \neq h(x)]$


where  $\Pr_{x \in D}$  excludes objects used in training  $h$ .

# PAC error



- There is a (hidden) concept
- There is given training data (3 positive, 3 negative)

# PAC error



- Here the true error is non-zero, although  $h$  and  $c$  agree on all six training instances (training error = 0).
- How probable is it that the observed training error gives a misleading estimate of the true error?

# PAC learnable

- Characterise target concepts that can be reliably learned from (1) a 'reasonable' number of (randomly drawn) training examples and (2) a 'reasonable' amount of computation.

# PAC learnable

- Characterise target concepts that can be reliably learned from (1) a 'reasonable' number of (randomly drawn) training examples and (2) a 'reasonable' amount of computation.
- Sometimes we have an unlucky draw of examples
- With finite number of training examples there are hypotheses that work identical on the training examples: how to choose?
- We will not demand zero error, but an arbitrarily small error (approximately correct)
- We will not demand small error on all training sets, but that the failure is bounded (probably correct)

## Probably Approximately Correct (PAC)




# PAC learnable

- Characterize target concepts that can be reliably learned from (1) a 'reasonable' number of (randomly drawn) training examples and (2) a 'reasonable' amount of computation.
- We will not demand small error on all training sets, but that the failure is bounded (probably correct)

## Probably Approximately Correct (PAC)


- $C$  is PAC-learnable by  $L$  using  $H$  if for all  $c$ , distribution  $D$ , the learner  $L$  will with probability at least  $(1 - \delta)$  output a hypothesis  $h$  such that  $\text{error}_D(h) \leq \varepsilon$  in time/samples that is polynomial in  $1/\varepsilon$ ,  $1/\delta$ ,  $\text{size}(c)$ , dimensionality  $n$

# Example: rectangle learning




- Learn an axis-parallel rectangle  $R$  from  $+$  and  $-$  examples in  $\mathbb{R}^2$
- Examples are randomly drawn from  $D$
- Adapt hypothesis rectangle  $R'$  to approximate  $R$

# Example: rectangle learning




- The error of  $R'$  is  $(R - R') \cup (R' - R)$
- What learning strategy to use so we can efficiently learn it?...

# Example: rectangle learning




- Use the '**tightest** fit rectangle' (definition of  $L$ !):  $R'$
- We make still an error on the test set:  $R'$  is always contained in  $R$
- Can we analyse the error? We can split the error in four strips (shown one in orange:  $T'$ ).

# Example: rectangle learning



- What is the prob. the learner has error larger than  $\varepsilon$ ?
- Now define a new strip,  $T$
- Strip  $T$  is 'grown' such that it covers  $\varepsilon/4$  of the prob.mass (for given  $\varepsilon$ )
- Now  $T$  may cover  $T'$  or may not cover  $T'$


# Example: rectangle learning



- If  $T$  covers  $T'$  and that holds for all strips, then the error


$$\begin{aligned} P[\text{error}] &= P[\text{yellow}] \leq P[T_1] + P[T_2] + P[T_3] + P[T_4] \\ &= 4(\varepsilon/4) = \varepsilon \end{aligned}$$

# Example: rectangle learning



- Can we now estimate the probability that  $T$  does not cover  $T'$  (that the error exceeds  $\varepsilon$ )?
- Can we show that, with sufficient number of training samples,  $R'$  will always be so large that  $T$  covers  $T'$ ? And how many training samples then?

# Example: rectangle learning



- T would have covered T' when none of the positive samples would have hit area T


$$P[\text{random } x \text{ hits } T] = \varepsilon/4$$

$$P[\text{random } x \text{ missed } T] = 1 - \varepsilon/4$$

$$P[m \text{ random } x' \text{'s miss } T] = (1 - \varepsilon/4)^m$$



# Example: rectangle learning



- We have 4 strips, so

$$P[m \text{ random } x' \text{'s miss any } T\text{'s}] = P[\text{miss } T_1] \text{ or } P[\text{miss } T_2] \text{ or } \dots$$

- So, the probability that our  $R'$  has an error larger than  $\varepsilon$  is something we want to bound:

$$P[R' \text{ has larger error than } \varepsilon] \leq 4(1 - \varepsilon/4)^m < \delta$$

# Union bound

- For a countable number of events  $A_1, A_2, \dots$  the probability that **at least one** of the events happens, is no greater than the sum of the probabilities of the individual events:

$$P\left(\bigcup_i A_i\right) = P(A_1 \text{ or } A_2 \text{ or } \dots) \leq \sum_i P(A_i)$$

- If the events are mutually exclusive, then the equality holds

# Rectangle learning

- Now we want to bound the chance that our  $R'$  makes an error larger than  $\varepsilon$  by  $\delta$

$$4(1 - \varepsilon/4)^m < \delta$$

- Now use:  $e^{-x} \geq (1 - x)$

and we obtain:  $4e^{-m\varepsilon/4} \geq 4(1 - \varepsilon/4)^m$


- So instead we can demand:  $4e^{-m\varepsilon/4} < \delta$

$$-m\varepsilon/4 < \log(\delta/4)$$

$$m\varepsilon/4 > \log(4/\delta)$$


- This  $R'$  is PAC learnable!  $m > (4/\varepsilon) \log(4/\delta)$

# More 'interpretation'...




- When I get a few training samples
- then the true error of  $R'$  may still be anything
- ... in particular when  $m$  is small.

# More interpretation...



- When I get more samples, my error tends to become smaller
- But still, I may be unlucky

# More interpretation...



- But still, I may be unlucky: for all errors I still have some probability  $\delta$  that my classifier  $R'$  actually is worse than that (red area)

# 'Conclusion'


- So the general question in Learning Theory is: How many samples  $m$  do I need such that my learner  $L$  gives a classifier with small error?
- Is the number of samples  $m$  'reasonable' (i.e. not too large)?

**Special case:  
Discrete Hypothesis spaces  
and  
Consistent learners**



# Consistent learner


- The Version Space is the collection of all consistent hypotheses (zero error on training set  $S$ , test error can be anything)



- Consistent learner has zero error on training set

# $\varepsilon$ -exhausted version space

- The Version Space is the collection of all consistent hypotheses (zero error on training set  $S$ , test error can be anything)



- In an  $\varepsilon$ -exhausted version space, all hypotheses have an error smaller than  $\varepsilon$  on the test set. GOOD!

# When not exhausted?

- What is the probability that the version space is not  $\varepsilon$ -exhausted?
- Assume there are  $k$  hypotheses with error larger than  $\varepsilon$
- We fail to exhaust the version space if any of these hypotheses is consistent with our training sample (with  $m$  training objects)
- The probability that a hypothesis with error larger than  $\varepsilon$  is consistent with  $m$  objects is at most  $(1 - \varepsilon)^m$

# The probability is...

- Given  $k$  hypotheses with error  $> \varepsilon$ , the probability that at least one of them is consistent with all  $m$  training examples is at most:  $k(1 - \varepsilon)^m$  (union bound)

- Obviously  $k \leq |H|$  and using  $(1 - x) \leq e^{-x}$

$$k(1 - \varepsilon)^m \leq |H|(1 - \varepsilon)^m \leq |H|e^{-\varepsilon m}$$

# Which means...

- Given  $k$  hypotheses with error  $> \varepsilon$ , the probability that at least one of them is consistent with all  $m$  training examples is at most:  $k(1 - \varepsilon)^m$

- Obviously  $k \leq |H|$  and using  $(1 - x) \leq e^{-x}$

$$k(1 - \varepsilon)^m \leq |H|(1 - \varepsilon)^m \leq |H|e^{-\varepsilon m}$$

- So when we want to bound the chance of having a failure:  $|H|e^{-\varepsilon m} \leq \delta$

we need:  $m \geq \frac{1}{\varepsilon} (\ln |H| + \ln(1/\delta))$

# Consistent learners

- We found a very general bound for ANY consistent learner: 
$$m \geq \frac{1}{\varepsilon} (\ln |H| + \ln(1/\delta))$$
- It depends on the (log of the) size of the hypothesis space
- This number  $m$  of training examples is sufficient to assure that any consistent hypothesis will be probably (with prob.  $(1 - \delta)$ ) approximately (within error  $\varepsilon$ ) correct.
- Note that we assumed consistent algorithms: zero training error in a discrete feature space...

# VC-dimension

- We discussed discrete feature spaces and hypothesis spaces with zero class overlap (the learner can perfectly learn the concept)
  - Inconsistent learners are also possible (weak learners, later in lecture): the bounds gets less tight
  - More class overlap is possible, but too much for this lecture...
- 
- What if we use **continuous** feature/hypotheses spaces?
  - We have seen it in the Pattern Recognition course: Vapnik-Chervonenkis dimension

# Bounding the true error (Copied from the PR course)

With probability at least  $1 - \eta$  the inequality holds:

$$\varepsilon \leq \varepsilon_A + \frac{\mathcal{E}(N)}{2} \left( 1 + \sqrt{1 + \frac{\varepsilon_A}{\mathcal{E}(N)}} \right)$$

where

$$\mathcal{E}(N) = 4 \frac{h(\ln(2N/h) + 1) - \ln(\eta/4)}{N}$$

V.Vapnik, Statistical learning theory, 1998

- When  $h$  is small, the true error is close to the apparent error



# VC-dimension and samples

- When you have the VC-dimension of a learner  $L$ , then holds:

$$m \geq c_0 \left( \underbrace{\frac{1}{\varepsilon} \log \frac{1}{\delta}}_{\text{(similar to the discrete feature space)}} + \underbrace{\frac{h}{\varepsilon} \log \frac{1}{\varepsilon}}_{\text{(caused by the continuous feature space)}} \right)$$

- This VC-dimension is the analogue of  $|H|$
- Similarly, also lower bounds on the number of training samples can be given.
- Only bounds/approximations on the VC-dimension are known for most classifiers

# 'No Free Lunch Theorem'

- There are **no** context-independent reasons to favor one learner over another (Wolpert, 1996)
- Averaged over all possible problems, all learners have the same average performance
- So when one algorithm seems to outperform another, it just fits the problem better, and it does not mean that the one algorithm is inherently better
- Claims in literature that this procedure/algorithm performs 'best' overall should be considered with some care...

# No free lunch...

- Assume we have a discrete feature space (with size  $|X|$ ), then we need for a consistent learner that the number of samples  $m$ :

$$m \geq \frac{1}{\varepsilon} (\ln |H| + \ln(1/\delta))$$

- But assume then that **ALL** possible hypotheses are allowed:

$$|H| = 2^{|X|}$$

- For a discrete binary feature space with  $n$  features:

$$|X| = 2^n$$

- The number of training examples grows exponentially:  $m \geq \frac{1}{\varepsilon} (2^n \ln 2 + \ln(1/\delta))$  **NOT PAC learnable**

# Conclusions

- General statements can be made about the number of required training objects, but additional (strong) assumptions on the feature space or hypothesis space have to be made
- Bounds can be given, but are often very loose (and not always easy to interpret)
- Sometimes constructive algorithms are invented (AdaBoost, Support Vector Machines)
- Averaged over all problems, all methods are equally good ... ..

# Weak/strong learners

- PAC learning requires that the error  $\varepsilon$  can be arbitrarily small, and the confidence  $1 - \delta$  can be set arbitrarily high.
- What if we have a **weak** learner that has a **fixed** error  $\varepsilon_0$  and confidence  $1 - \delta_0$ ?
- Magically, it appears that there is an algorithm that can use the weak learner to boost it to a full PAC learner (a **strong** learner)
- It also means that PAC learning is very general: the demands on the learner do not have to be that strict (you can always boost it)

# Original boosting

- The original idea: resample original data  $D$  such that more difficult objects are used to train 2 additional (weak) learners
- The collection of (3) weak learners predict the final label by majority voting
- Later versions do not resample the training set, but split the feature space, or introduce other combinations of weak learners

# AdaBoost

- Inspired by boosting a weak classifier to a strong one: Adaptive Boosting
- My explanation starts from assumptions on (1) the model, and (2) the error function. The (PAC) theory is not needed in the derivation.

- Assumption 1: the model is linear additive:

$$F_K(\mathbf{x}) = \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}) + \alpha_K f_K(\mathbf{x})$$

where

$$f_i(\mathbf{x}) = \pm 1$$

and  $\alpha_i$  are weights.

(binary outputs!)


# AdaBoost

- Assumption 2: the loss/error on a training set is measured by:

$$L = \sum_{i=1}^N \exp(-y_i F_K(\mathbf{x}_i))$$

exponential loss

0-1 loss





# AdaBoost

- To optimize both the weak classifiers  $f_i(\mathbf{x})$  and the weights  $\alpha_i$  is an open problem
- Instead, do it incrementally:

$$F_K(\mathbf{x}) = \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}) + \alpha_K f_K(\mathbf{x})$$

# AdaBoost

- To optimize both the weak classifiers  $f_i(\mathbf{x})$  and the weights  $\alpha_i$  is an open problem
- Instead, do it incrementally:

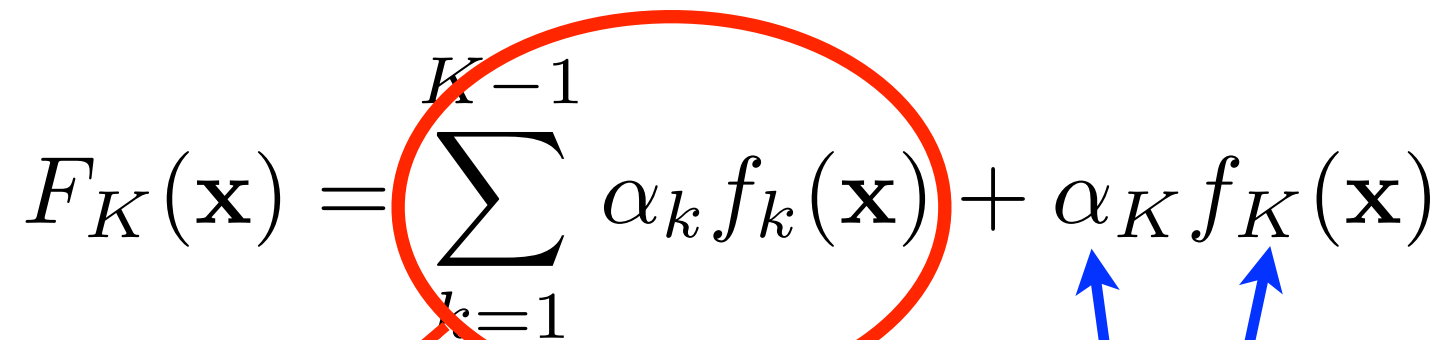
$$F_K(\mathbf{x}) = \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}) + \alpha_K f_K(\mathbf{x})$$

fix this

optimize these two terms

# AdaBoost

- To optimize both the weak classifiers  $f_i(\mathbf{x})$  and the weights  $\alpha_i$  is an open problem
- Instead, do it incrementally:

$$F_K(\mathbf{x}) = \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}) + \alpha_K f_K(\mathbf{x})$$


fix this

optimize these  
two terms

- Minimize L:

$$L = \sum_{i=1}^N \exp \left( -y_i \left[ \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}_i) + \alpha_K f_K(\mathbf{x}_i) \right] \right)$$

# AdaBoost

$$\begin{aligned} L &= \sum_{i=1}^N \exp \left( -y_i \left[ \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}_i) + \alpha_K f_K(\mathbf{x}_i) \right] \right) \\ &= \sum_{i=1}^N \underbrace{\exp \left( -y_i \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}_i) \right)}_{w_i} \exp \left( -y_i \alpha_K f_K(\mathbf{x}_i) \right) \\ &= \sum_{i=1}^N w_i \exp \left( -y_i \alpha_K f_K(\mathbf{x}_i) \right) \end{aligned}$$

- Now distinguish correctly and incorrectly classified objects:

$$y_i f_K(\mathbf{x}_i) = 1 \quad \rightarrow \quad \mathbf{x}_i \in C_K \quad (\text{correct})$$

$$y_i f_K(\mathbf{x}_i) = -1 \quad \rightarrow \quad \mathbf{x}_i \in W_K \quad (\text{wrong})$$

# AdaBoost

$$\begin{aligned} L &= \sum_{i=1}^N w_i \exp(-\alpha_K y_i f_K(\mathbf{x}_i)) && \text{(definition)} \\ &= \sum_{C_K} w_i \exp(-\alpha_K) + \sum_{W_K} w_i \exp(\alpha_K) && \text{(previous page)} \\ &= \sum_{i=1}^N w_i \exp(-\alpha_K) - \sum_{W_K} w_i \exp(-\alpha_K) + \sum_{W_K} w_i \exp(\alpha_K) && \text{(sum over correct = sum over N - sum over wrongs)} \\ &= \sum_{i=1}^N w_i \exp(-\alpha_K) + \sum_{W_K} w_i (\exp(\alpha_K) - \exp(-\alpha_K)) && \text{(rearrange)} \\ &= \sum_{i=1}^N w_i \exp(-\alpha_K) + \sum_{i=1}^N w_i (\exp(\alpha_K) - \exp(-\alpha_K)) \mathcal{I}(f_K(\mathbf{x}_i) \neq y_i) && \text{(replace sum over wrongs by sum over N)} \end{aligned}$$

# AdaBoost

- To minimize w.r.t.  $f_K$

$$L = \sum_{i=1}^N w_i \exp(-\alpha_K) + \sum_{i=1}^N w_i (\exp(\alpha_K) - \exp(-\alpha_K)) \mathcal{I}(f_K(\mathbf{x}_i) \neq y_i)$$

we should minimize  $\varepsilon_K = \sum_{i=1}^N w_i \mathcal{I}(f_K(\mathbf{x}_i) \neq y_i)$

- Or, in other words, we should find a classifier  $f_K$  that minimizes the error where each object is re-weighted by:

$$w_i = \exp \left( -y_i \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}_i) \right)$$

(how bad was  $\mathbf{x}_i$  classified by the previous  $F_{K-1}$ )

# AdaBoost

- Ok, so the classifier should minimize the weighted error, what about the weight  $\alpha_K$ ?
- Take derivative of the loss with respect to  $\alpha_K$  and set it to zero:

$$\frac{\partial L}{\partial \alpha_K} = -\exp(-\alpha_K) \sum_{i=1}^N w_i + (\exp(\alpha_K) + \exp(-\alpha_K)) \varepsilon_K = 0$$

where:

$$\varepsilon_K = \sum_{i=1}^N w_i \mathcal{I}(f_K(\mathbf{x}_i) \neq y_i)$$

- Solving it:

$$\sum_{i=1}^N w_i = (\exp(2\alpha_K) + 1) \varepsilon_K$$
$$\alpha_K = \frac{1}{2} \log \left( \frac{\sum_i w_i}{\varepsilon_K} - 1 \right)$$

# AdaBoost

1. Give each object a weight  $w_i = 1$
2. Train a classifier that minimizes the weighted error:

$$\varepsilon_K = \sum_{i=1}^N w_i \mathcal{I}(f_K(\mathbf{x}_i) \neq y_i)$$

3. Compute the weight of the classifier:

$$\alpha_K = \frac{1}{2} \log \left( \frac{\sum_i w_i}{\varepsilon_K} - 1 \right)$$

4. Compute the new object weights:

$$w_i = \exp \left( -y_i \sum_{k=1}^K \alpha_k f_k(\mathbf{x}_i) \right)$$

5. If K not large enough, go to 2, else we're done:



$$F_K(\mathbf{x}) = \sum_{k=1}^K \alpha_k f_k(\mathbf{x})$$



# AdaBoost


- Use a simple decision stump for weak classifier
- Compute  $\alpha_K$  and reweigh each object using

$$w_i = \exp \left( -y_i \sum_{k=1}^{K-1} \alpha_k f_k(\mathbf{x}_i) \right)$$




# AdaBoost

- Train a new decision stump on reweighted objects
- Recompute  $\alpha_K$  and  $w_i$
- Repeat and repeat...



# AdaBoost

- Finally, we end up with:



$$F_K(\mathbf{x}) = \sum_{k=1}^K \alpha_k f_k(\mathbf{x})$$

# Conclusions

- PAC learning: theoretical bounds on errors, required sample sizes
- Bounds are hard to get, and not very tight
- Useful by-product: AdaBoost